



Mahavir Education Trust's
SHAH & ANCHOR KUTCHHI ENGINEERING COLLEGE
Chembur, Mumbai - 400 088
UG Program in Artificial Intelligence and Data Science

EXPERIMENT 2

AIM: Simple Linear Regression in Python/R.

THEORY :

- Linear Regression is a machine learning algorithm based on supervised learning.
- Linear regression is used for finding linear relationship between target and one or more predictors.
- Linear regression is one of the easiest and most popular Machine Learning algorithms.
- Linear regression algorithm shows a linear relationship between a dependent (y) and one or more independent (x) variables, hence called as linear regression.
- Since linear regression shows the linear relationship, which means it finds how the value of the dependent variable is changing according to the value of the independent variable.
- It is a statistical method that is used for predictive analysis.
- Linear regression makes predictions for continuous/real or numeric variables such as sales, salary, age, product price, etc.
- There are two types of linear regression- Simple Regression and Multiple Regression.

❖ **SIMPLE LINEAR REGRESSION:**

- Simple linear regression is a statistical method for establishing the relationship between two variables using a straight line.
- The line is drawn by finding the slope and intercept, which define the line and minimize regression errors.
- The simplest form of simple linear regression has only one x variable and one y variable.
- The x variable is the independent variable because it is independent of what you try to predict the dependent variable.
- The y variable is the dependent variable because it depends on what you try to predict.
- $y = B_0 + B_1x + e$ is the formula used for simple linear regression.
- y is the predicted value of the dependent variable (y) for any given value of the independent variable (x).



Mahavir Education Trust's
SHAH & ANCHOR KUTCHHI ENGINEERING COLLEGE
Chembur, Mumbai - 400 088
UG Program in Artificial Intelligence and Data Science

- B_0 is the intercept, the predicted value of y when the x is 0.
- B_1 is the regression coefficient – how much we expect y to change as x increases.
- x is the independent variable (the variable we expect is influencing y).
- e is the error of the estimate, or how much variation there is in our regression coefficient estimate.
- Simple linear regression establishes a line that fits your data, but it does not guarantee that the line is good enough.
- For example, if your data points have an upward trend and are very far apart, then simple linear regression will give you a downward-sloping line, which will not match your data.
- A simple linear regression can accurately capture the relationship between two variables in simple relationships.
- But when dealing with more complex interactions that require more thought, you need to switch from simple to multiple regression.

CODE OUTPUT (PYTHON)

Dataset used:- Student_Score

```
colab.research.google.com/drive/1CO-dYmbqdBj7ni259jKSL3731zg1c3bD#scrollTo=OibDCzQdMsie
```

+ Code + Text

```
[1] #Import Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import math
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
df = pd.read_csv("student_scores.csv")
df.head()
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

```
[4] x = df["Hours"]
y = df["Scores"]
```

```
[5] df.describe()
```



Mahavir Education Trust's
SHAH & ANCHOR KUTCHHI ENGINEERING COLLEGE
Chembur, Mumbai - 400 088
UG Program in Artificial Intelligence and Data Science

colab.research.google.com/drive/1CO-dYmbqdBj7ni259jKSL3731zg1c3bD#scrollTo=OibDCzQdMsie

+ Code + Text

RAM Disk

min 1.100000 17.000000
25% 2.700000 30.000000
50% 4.800000 47.000000
75% 7.400000 75.000000
max 9.200000 95.000000

df.dtypes

Hours float64
Scores int64
dtype: object

[7] df.info()

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 25 entries, 0 to 24  
Data columns (total 2 columns):  
#   Column  Non-Null Count  Dtype  
--  --  
0   Hours    25 non-null      float64  
1   Scores   25 non-null      int64  
dtypes: float64(1), int64(1)  
memory usage: 528.0 bytes
```

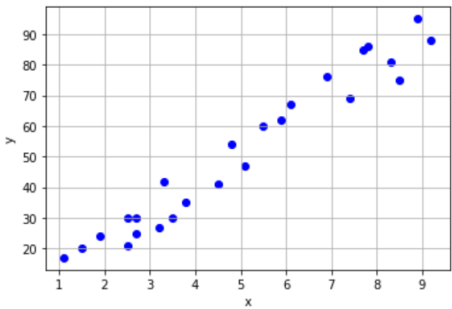
[8] plt.scatter(x, y, color = 'Blue')
#plt.plot(x, y_predicted, color = 'green')
plt.xlabel('x')

colab.research.google.com/drive/1CO-dYmbqdBj7ni259jKSL3731zg1c3bD#scrollTo=OibDCzQdMsie

+ Code + Text

RAM Disk

[8] plt.scatter(x, y, color = 'Blue')
#plt.plot(x, y_predicted, color = 'green')
plt.xlabel('x')
plt.ylabel('y')
plt.grid()
plt.show()



[9] df=df.drop_duplicates()

[10] x.head()

0 2.5
- -



Mahavir Education Trust's
SHAH & ANCHOR KUTCHHI ENGINEERING COLLEGE
Chembur, Mumbai - 400 088
UG Program in Artificial Intelligence and Data Science

+ Code + Text

✓ [10] x.head()
0s

```
0    2.5
1    5.1
2    3.2
3    8.5
4    3.5
Name: Hours, dtype: float64
```

✓ [11]
0s

```
x = np.array(x)
y = np.array(y)
y

array([21, 47, 27, 75, 30, 20, 88, 60, 81, 25, 85, 62, 41, 42, 17, 95, 30,
       24, 67, 69, 30, 54, 35, 76, 86])
```

✓ [12] x = x.reshape(-1,1)
0s
y = y.reshape(-1,1)

✓ [13]
0s

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2)

print(len(x_test))
```

✓ 0s completed at 3:34 PM




Mahavir Education Trust's
SHAH & ANCHOR KUTCHHI ENGINEERING COLLEGE
Chembur, Mumbai - 400 088
UG Program in Artificial Intelligence and Data Science

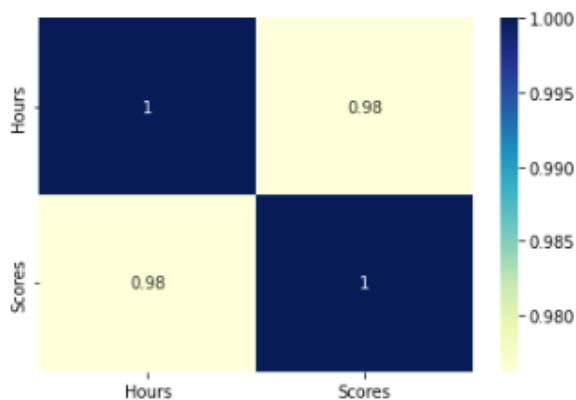
+ Code + Text

```
✓ [22] import seaborn as sns  
1s     sns.regplot(x, y, ci=None)
```

```
/usr/local/lib/python3.8/dist-packages/seaborn/_decorators.py:36: FutureWarning:  
warnings.warn(  
<matplotlib.axes._subplots.AxesSubplot at 0x7f71fcbac100>
```



```
✓ [24] sns.heatmap(df.corr(), cmap="YlGnBu", annot=True)  
1s     plt.show()
```



```
✓ [25] import statsmodels.api as sm  
1s
```

```
✓ [26] x_train_sm=sm.add_constant(x_train)  
0s     lr=sm.OLS(y_train,x_train_sm).fit()
```

```
✓ [27] lr.summary()  
0s
```



OLS Regression Results

Dep. Variable:	y	R-squared:	0.947
Model:	OLS	Adj. R-squared:	0.944
Method:	Least Squares	F-statistic:	322.4



Mahavir Education Trust's
SHAH & ANCHOR KUTCHHI ENGINEERING COLLEGE
Chembur, Mumbai - 400 088
UG Program in Artificial Intelligence and Data Science

✓ [27] lr.summary()
0s

```
OLS Regression Results

Dep. Variable:  y                R-squared:    0.947
Model:         OLS              Adj. R-squared: 0.944
Method:        Least Squares    F-statistic: 322.4
Date:          Tue, 14 Feb 2023  Prob (F-statistic): 6.14e-13
Time:          09:25:52         Log-Likelihood: -62.677
No. Observations: 20            AIC:         129.4
Df Residuals:    18            BIC:         131.3
Df Model:        1
Covariance Type: nonrobust

   coef  std err   t    P>|t| [0.025 0.975]
const 1.9738 3.195   0.618  0.544 -4.738 8.686
x1     9.7897 0.545  17.957 0.000  8.644 10.935

Omnibus:   7.470   Durbin-Watson:  2.620
Prob(Omnibus): 0.024   Jarque-Bera (JB): 1.836
Skew:      -0.090     Prob(JB):   0.399
Kurtosis:   1.527     Cond. No.   14.6
```

Code and Output: R

```
data = read.csv("C:/Users/saksh/Downloads/student_scores.csv")
library(dplyr)
# random_Inspection of Dataset
sample_n(data,5)
```


```
> data = read.csv("C:/Users/saksh/Downloads/student_scores.csv")
> library(dplyr)

  Hours Scores
1    4.5     41
2    4.8     54
3    8.5     75
4    7.8     86
5    6.9     76
> plot(data$Hours,data$score)
> |
```

```
# plotting data to understand which regression method should be applied
plot(data$Hours,data$Score)
```



Mahavir Education Trust's
SHAH & ANCHOR KUTCHHI ENGINEERING COLLEGE
Chembur, Mumbai - 400 088
UG Program in Artificial Intelligence and Data Science



Pearson's correlation test will also help to check linearity and correlation between data
`cor.test(data$Hours,data$Score)`

#Since, p-value is less than 0.05 and sample correlation is 0.9782416.

#Hence, High significant correlation (Highly Linearly Related) exists in the population.

`library(caret)`

```
> sample_n(data,5)
  Hours Scores
1   4.5    41
2   4.8    54
3   8.5    75
4   7.8    86
5   6.9    76
> plot(data$Hours,data$Score)
> cor.test(data$Hours,data$Score)

Pearson's product-moment correlation

data: data$Hours and data$Score
t = 21.583, df = 23, p-value < 2.2e-16
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.9459248 0.9896072
sample estimates:
      cor 
0.9761907
```

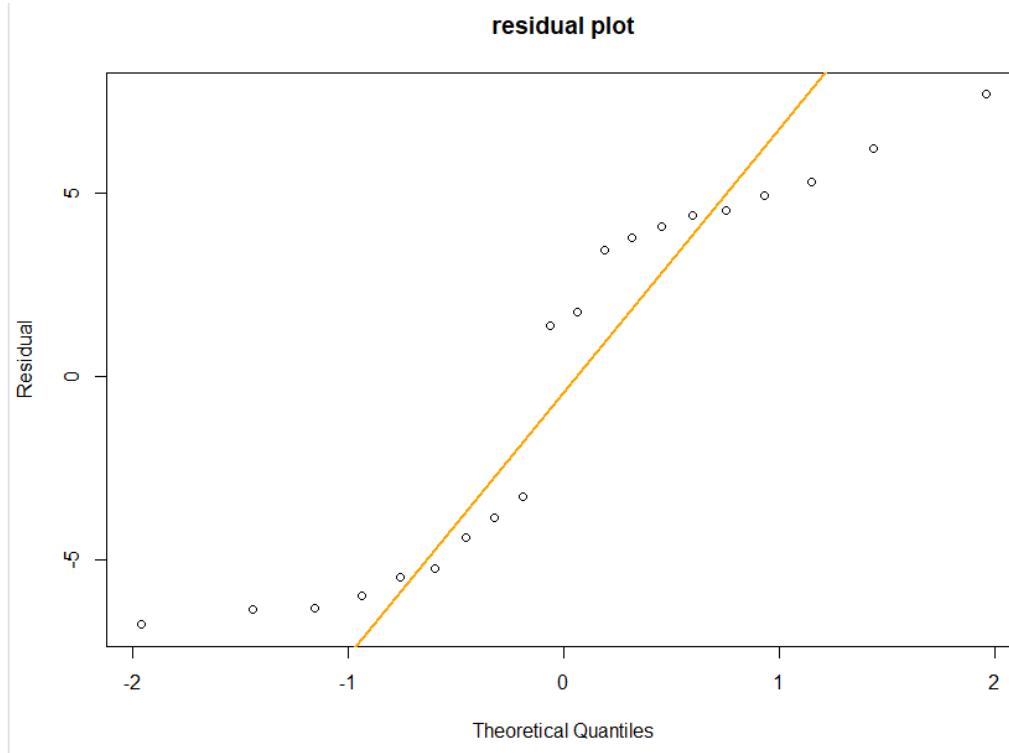
#Split the given dataset into train and test data,

#fit the model and obtain summary of model as follows



Mahavir Education Trust's
SHAH & ANCHOR KUTCHHI ENGINEERING COLLEGE
Chembur, Mumbai - 400 088
UG Program in Artificial Intelligence and Data Science


```
training.samples <- data$Score %>%  
  createDataPartition(p = 0.7, list = FALSE)  
train.data <- data[training.samples, ]  
test.data <- data[-training.samples, ]  
  
# Fit model  
model <- lm(Scores ~ Hours, data = train.data)  
summary(model)  
  
# Visualization  
qqnorm(model$residuals, ylab="Residual", main="residual plot")  
qqline(model$residuals, col = "orange", lwd = 2)
```



```
# Making prediction  
prediction <- model %>% predict(test.data)  
# Visualization  
plot(test.data$Scores , prediction)  
abline(lm(prediction ~ Scores, data = test.data), col = "blue")
```




Mahavir Education Trust's
SHAH & ANCHOR KUTCHHI ENGINEERING COLLEGE
Chembur, Mumbai - 400 088
UG Program in Artificial Intelligence and Data Science



```
# Statistical Measure
data.frame( R2 = R2(prediction, test.data$Scores),
            RMSE = RMSE(prediction, test.data$Scores),
            MAE = MAE(prediction, test.data$Scores))

# Multifold training
# Define training control
train.control <- trainControl(method = "repeatedcv",
                              number = 4, repeats = 3)

# Train the model
model_cv <- train(Scores ~ Hours , data = data, method="lm",
                  trControl = train.control)

# Summarize the results
print(model_cv)
```



Mahavir Education Trust's
SHAH & ANCHOR KUTCHHI ENGINEERING COLLEGE
Chembur, Mumbai - 400 088
UG Program in Artificial Intelligence and Data Science

```
> training.samples <- data$Score %>%
+   createDataPartition(p = 0.7, list = FALSE)
> train.data <- data[training.samples, ]
> test.data <- data[-training.samples, ]
> model <- lm(Scores ~ Hours, data = train.data)
> summary(model)

Call:
lm(formula = Scores ~ Hours, data = train.data)

Residuals:
    Min       1Q   Median       3Q      Max
-6.755 -5.282  1.579   4.420   7.686

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)   0.9591     2.7716   0.346   0.733
Hours         10.1075     0.4988  20.266 7.67e-14 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.269 on 18 degrees of freedom
Multiple R-squared:  0.958,    Adjusted R-squared:  0.9557
F-statistic: 410.7 on 1 and 18 DF,  p-value: 7.67e-14

> qqnorm(model$residuals, ylab="Residual", main="residual plot")
> qqline(model$residuals, col = "orange", lwd = 2)
> prediction <- model %>% predict(test.data)
> plot(test.data$Scores, prediction)
> abline(lm(prediction ~ Scores, data = test.data), col = "blue")
> # Statistical Measure
> data.frame( R2 = R2(prediction, test.data$Scores),
+             RMSE = RMSE(prediction, test.data$Scores),
+             MAE = MAE(prediction, test.data$Scores))
  R2      RMSE      MAE
1 0.9492178 6.923496 6.261933
> # Multifold training
> # Define training control
> train.control <- trainControl(method = "repeatedcv",
+                               number = 4, repeats = 3)
> # Train the model
> model_cv <- train(Scores ~ Hours, data = data, method="lm",
+                  trControl = train.control)
> # Summarize the results
> print(model_cv)
Linear Regression

25 samples
1 predictor

No pre-processing
Resampling: Cross-validated (4 fold, repeated 3 times)
Summary of sample sizes: 18, 20, 20, 17, 19, 19, ...
Resampling results:
```



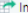


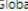






```
+                               trControl = train.control)
> # Summarize the results
> print(model_cv)
Linear Regression

25 samples
1 predictor

No pre-processing
Resampling: Cross-validated (4 fold, repeated 3 times)
Summary of sample sizes: 18, 20, 20, 17, 19, 19, ...
Resampling results:

RMSE      Rsquared  MAE
5.845614  0.9641844  5.425861

Tuning parameter 'intercept' was held constant at a value of TRUE
> |
```

Environment	History	Connections	Tutorial
 	 Import Dataset ▾	 388 MiB ▾	
R ▾  Global Environment ▾			
Data			
 data	25 obs. of 2 variables		
 model	List of 12		
 model_cv	List of 24		
 test.data	5 obs. of 2 variables		
 train.control	List of 27		
 train.data	20 obs. of 2 variables		
training.samples	int [1:20, 1] 1 3 5 7 8 9 10 12 13 14 ...		
values			
prediction	Named num [1:5] 52.5 86.9 16.1 78.8 20.2		

Conclusion:

regression analysis is a supervised learning algorithm that uses labeled data to produce continuous variables. The linear regression model comprises a single parameter and a linear connection between the dependent and independent variables. we have studied and implemented the basic concepts of Simple Linear Regression of both Python & R, using students-score dataset.